

Testing

November 4, 2024

1 COMP1801 - Machine Learning Coursework Solution

Let's start by importing the essential Python libraries for data analysis and machine learning.

```
[23]: # Import libraries
try:
    import glob
    import pandas as pd
    import numpy as np
    import matplotlib.pyplot as plt
    import seaborn as sns

    # Importing libraries for model building
    from sklearn.preprocessing import StandardScaler
    from sklearn.linear_model import Ridge
    from sklearn.ensemble import RandomForestRegressor
    from sklearn.linear_model import LogisticRegression
    from sklearn.ensemble import RandomForestClassifier
    from sklearn.model_selection import cross_val_score

    from statsmodels.stats.outliers_influence import variance_inflation_factor

except Exception as e:
    print(f"Error : {e}")
```

```
[24]: # Find the CSV file in the Datasets directory
data_path = '../Datasets/*.csv'
file_list = glob.glob(data_path)

for file in file_list:
    print(f"Found file: {file}")

# Ensure there is exactly one file
if len(file_list) == 1:
    # Load the dataset
    df = pd.read_csv(file_list[0])
    print(f"Loaded dataset: {file_list[0]}")
else:
```

```
raise FileNotFoundError("No CSV file found or multiple CSV files found in_
↳the Datasets directory.")
```

```
Found file: ../Datasets/COMP1801_Coursework_Dataset.csv
Loaded dataset: ../Datasets/COMP1801_Coursework_Dataset.csv
```

```
[25]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 16 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Lifespan              1000 non-null   float64
1   partType              1000 non-null   object
2   microstructure        1000 non-null   object
3   coolingRate            1000 non-null   int64
4   quenchTime            1000 non-null   float64
5   forgeTime             1000 non-null   float64
6   HeatTreatTime         1000 non-null   float64
7   Nickel%               1000 non-null   float64
8   Iron%                 1000 non-null   float64
9   Cobalt%               1000 non-null   float64
10  Chromium%             1000 non-null   float64
11  smallDefects           1000 non-null   int64
12  largeDefects           1000 non-null   int64
13  sliverDefects          1000 non-null   int64
14  seedLocation           1000 non-null   object
15  castType               1000 non-null   object
dtypes: float64(8), int64(4), object(4)
memory usage: 125.1+ KB
```

```
[26]: # Check for missing values
df.isnull().sum()
```

```
[26]: Lifespan          0
partType            0
microstructure      0
coolingRate         0
quenchTime          0
forgeTime           0
HeatTreatTime       0
Nickel%             0
Iron%               0
Cobalt%             0
Chromium%           0
smallDefects        0
largeDefects        0
```

```

sliverDefects      0
seedLocation       0
castType           0
dtype: int64

```

```
[27]: df.head()
```

```

[27]:   Lifespan  partType microstructure  coolingRate  quenchTime  forgeTime  \
0    1469.17   Nozzle    equiGrain         13         3.84         6.47
1    1793.64   Block    singleGrain        19         2.62         3.48
2     700.60   Blade    equiGrain         28         0.76         1.34
3    1082.10   Nozzle    colGrain          9         2.01         2.19
4    1838.83   Blade    colGrain         16         4.13         3.87

      HeatTreatTime  Nickel%  Iron%  Cobalt%  Chromium%  smallDefects  \
0           46.87    65.73  16.52   16.82         0.93           10
1           44.70    54.22  35.38    6.14         4.26           19
2            9.54    51.83  35.95    8.81         3.41           35
3           20.29    57.03  23.33   16.86         2.78            0
4           16.13    59.62  27.37   11.45         1.56           10

      largeDefects  sliverDefects  seedLocation  castType
0                0              0      Bottom      Die
1                0              0      Bottom  Investment
2                3              0      Bottom  Investment
3                1              0          Top  Continuous
4                0              0          Top      Die

```

```
[28]: df.describe()
```

```

[28]:   Lifespan  coolingRate  quenchTime  forgeTime  HeatTreatTime  \
count  1000.000000  1000.000000  1000.000000  1000.000000  1000.000000
mean    1298.556320    17.639000    2.764230    5.464600    30.194510
std     340.071434     7.491783    1.316979    2.604513    16.889415
min     417.990000     5.000000    0.500000    1.030000     1.030000
25%    1047.257500    11.000000    1.640000    3.170000    16.185000
50%    1266.040000    18.000000    2.755000    5.475000    29.365000
75%    1563.050000    24.000000    3.970000    7.740000    44.955000
max    2134.530000    30.000000    4.990000   10.000000    59.910000

      Nickel%      Iron%      Cobalt%      Chromium%  smallDefects  \
count  1000.000000  1000.000000  1000.000000  1000.000000  1000.000000
mean     60.243080   24.553580   12.434690    2.768650   17.311000
std       5.790475    7.371737    4.333197    1.326496   12.268365
min      50.020000    6.660000    5.020000    0.510000    0.000000
25%      55.287500   19.387500    8.597500    1.590000    7.000000
50%      60.615000   24.690000   12.585000    2.865000   18.000000

```

75%	65.220000	29.882500	16.080000	3.922500	26.000000
max	69.950000	43.650000	19.990000	4.990000	61.000000

	largeDefects	sliverDefects
count	1000.000000	1000.000000
mean	0.550000	0.292000
std	1.163982	1.199239
min	0.000000	0.000000
25%	0.000000	0.000000
50%	0.000000	0.000000
75%	0.000000	0.000000
max	4.000000	8.000000

```
[29]: # Using nunique()
num_parts = df['partType'].nunique()
print(f"Number of unique parts types: {num_parts}")

# Or using value_counts() to see the distribution
parts_distribution = df['partType'].value_counts()
print("\nDistribution of parts types:")
print(parts_distribution)
```


Number of unique parts types: 4

Distribution of parts types:

```
partType
Valve      265
Block      253
Nozzle     245
Blade      237
Name: count, dtype: int64
```

1.1 Distribution of Lifespan

```
[30]: # Plotting the distribution of 'Lifespan'
plt.figure(figsize=(8, 6))
sns.histplot(df['Lifespan'], bins=30, kde=True)
plt.title('Distribution of Lifespan')
plt.xlabel('Lifespan (hours)')
plt.ylabel('Frequency')
plt.show()
```



1.2 Distribution of Lifespan Between Different Metal Parts

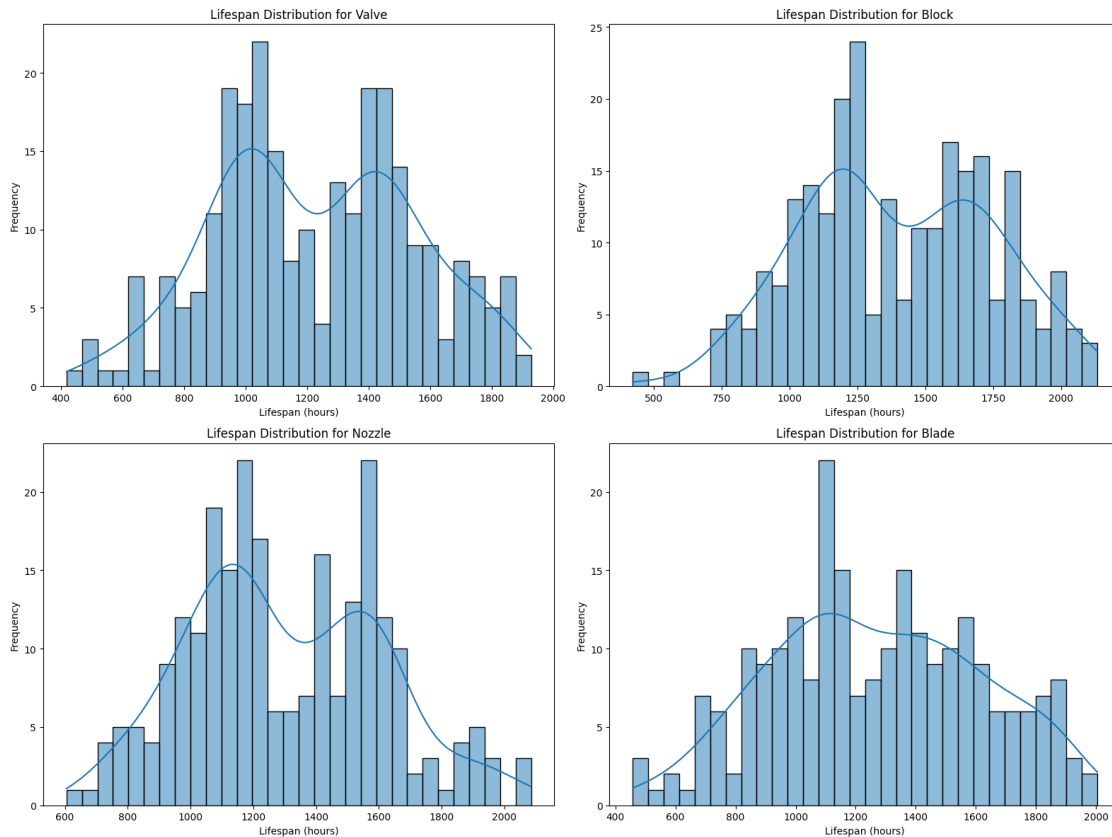
```
[31]: # List of metal parts
part_types = ['Valve', 'Block', 'Nozzle', 'Blade']

# Set up the subplots: 2 rows and 2 columns (one for each metal part)
fig, axes = plt.subplots(nrows=2, ncols=2, figsize=(16, 12))
axes = axes.flatten() # Flatten the axes array for easier iteration

# Loop through each metal part and create a histogram for its lifespan
# ↪ distribution
for i, part in enumerate(part_types):
    subset = df[df['partType'] == part]
    sns.histplot(subset['Lifespan'], bins=30, kde=True, ax=axes[i])
    axes[i].set_title(f'Lifespan Distribution for {part}')
    axes[i].set_xlabel('Lifespan (hours)')
    axes[i].set_ylabel('Frequency')

# Adjust layout for better visualization
```

```
plt.tight_layout()
plt.show()
```



1.3 Scatter Plot Analysis of Lifespan

```
[32]: # List of numerical columns excluding 'Lifespan'
numerical_cols = ['coolingRate', 'quenchTime', 'forgeTime', 'HeatTreatTime',
                  'Nickel%', 'Iron%', 'Cobalt%', 'Chromium%',
                  'smallDefects', 'largeDefects', 'sliverDefects']
```

```
[33]: # Determine the number of rows and columns for the subplots
num_cols = len(numerical_cols)
nrows = num_cols // 3 + (num_cols % 3 > 0) # Set up for 3 columns per row

# Create the subplots
fig, axes = plt.subplots(nrows=nrows, ncols=3, figsize=(20, nrows * 5))
axes = axes.flatten() # Flatten to easily iterate over

# Plot each scatterplot with a regression trend line in a different subplot
```

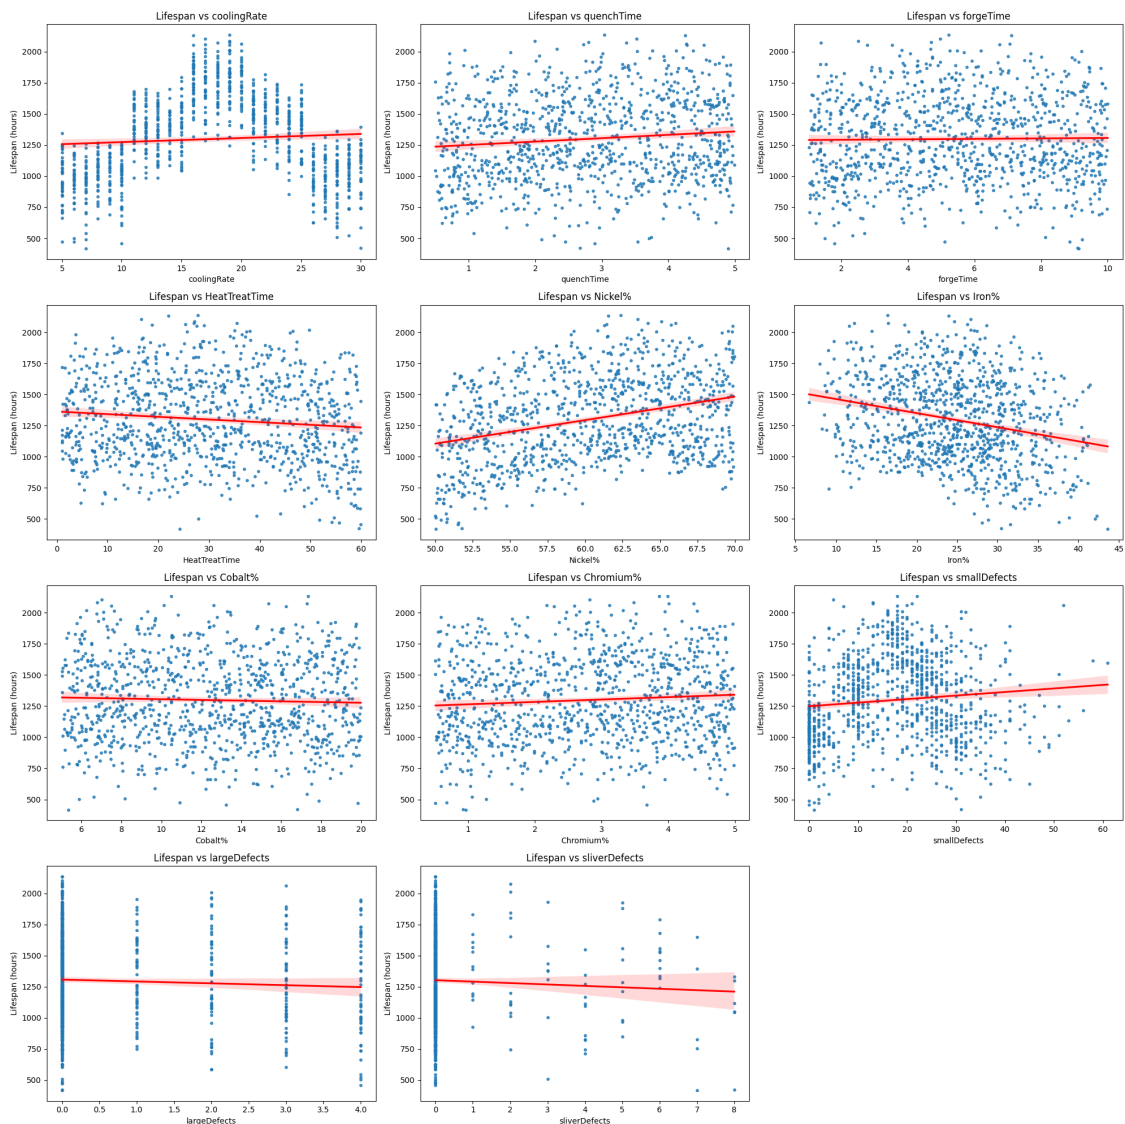
```

for i, col in enumerate(numerical_cols):
    sns.regplot(ax=axes[i], data=df, x=col, y='Lifespan', scatter_kws={'s': 10},
        line_kws={'color': 'red'})
    axes[i].set_title(f'Lifespan vs {col}')
    axes[i].set_xlabel(col)
    axes[i].set_ylabel('Lifespan (hours)')

# Remove any empty subplots if the number of features is not perfectly
# divisible by 3
for j in range(len(numerical_cols), len(axes)):
    fig.delaxes(axes[j])

# Adjust layout for better visualization
plt.tight_layout()
plt.show()

```



```
[38]: # List of categorical columns to include in the box plots
categorical_cols = ['partType', 'microstructure', 'seedLocation', 'castType']

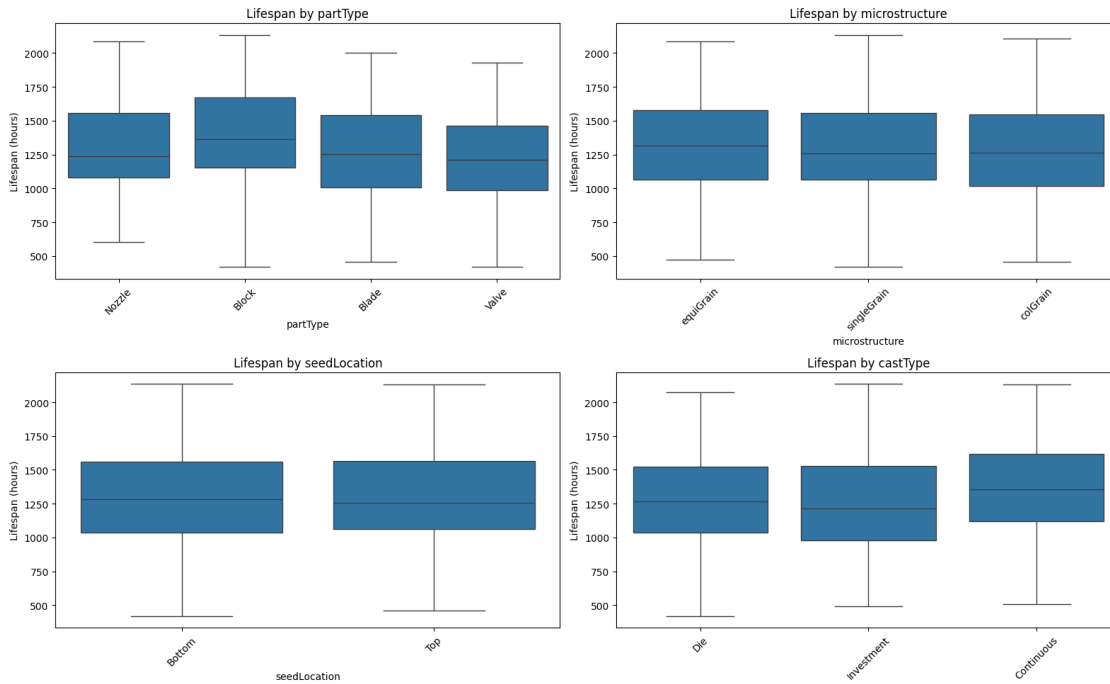
# Determine the number of rows and columns for the subplots
num_cols = len(categorical_cols)
nrows = num_cols // 2 + (num_cols % 2 > 0) # Set up for 2 columns per row for
↳ better readability

# Create the subplots
fig, axes = plt.subplots(nrows=nrows, ncols=2, figsize=(16, nrows * 5))
axes = axes.flatten() # Flatten the axes array to easily iterate over

# Plot each boxplot in a different subplot
for i, col in enumerate(categorical_cols):
    sns.boxplot(ax=axes[i], x=col, y='Lifespan', data=df)
    axes[i].set_title(f'Lifespan by {col}')
    axes[i].set_xlabel(col)
    axes[i].set_ylabel('Lifespan (hours)')
    axes[i].tick_params(axis='x', rotation=45)

# Remove any empty subplots if the number of features is odd
for j in range(len(categorical_cols), len(axes)):
    fig.delaxes(axes[j])

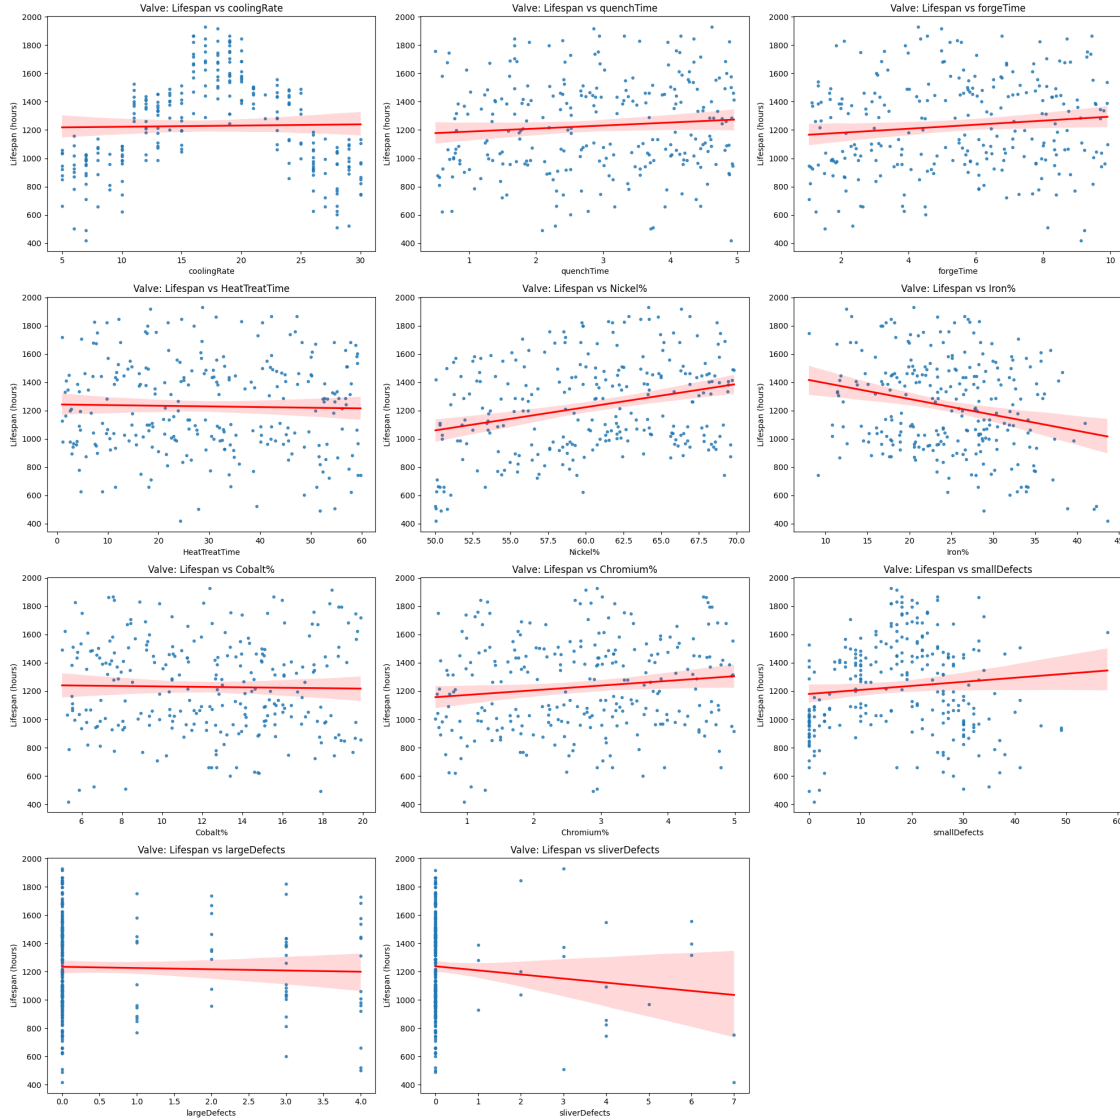
# Adjust layout for better visualization
plt.tight_layout()
plt.show()
```

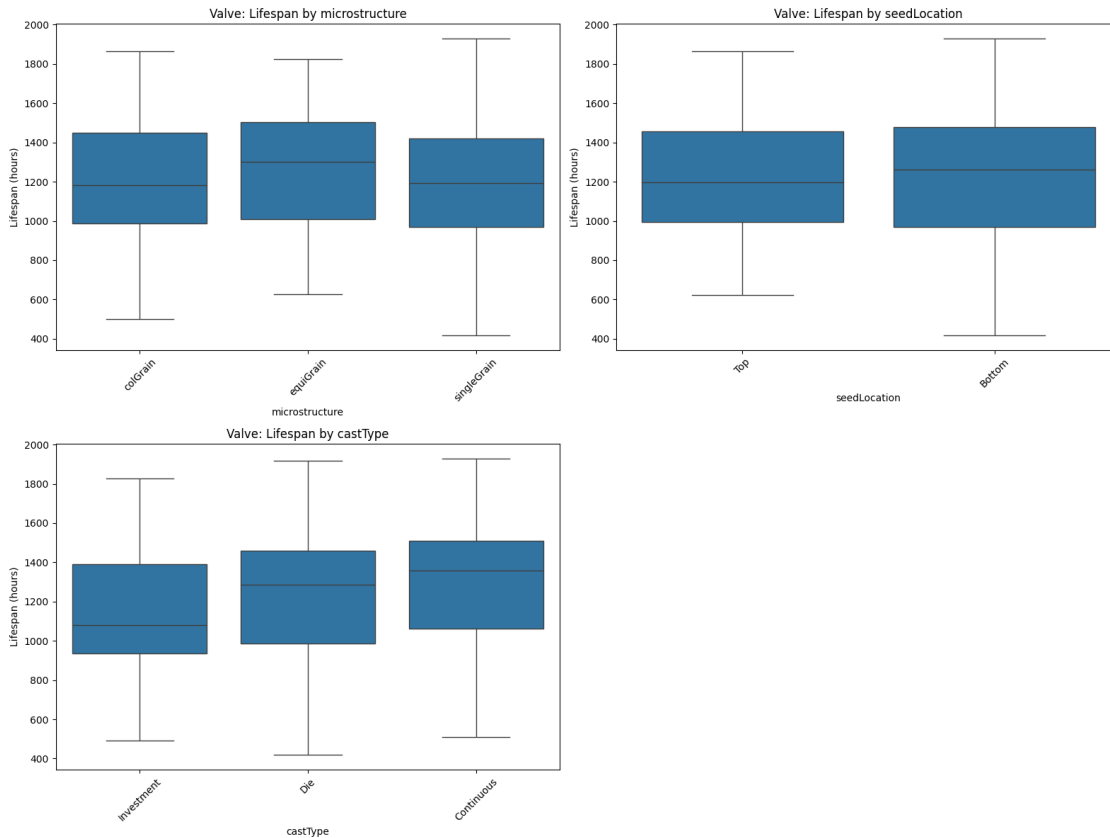
1.4 Analysing the data based on Part Type

1.4.1 Valve

```
[34]: # 1. Scatter Plots for Valve
valve_data = df[df['partType'] == 'Valve']
fig, axes = plt.subplots(nrows=4, ncols=3, figsize=(20, 20))
axes = axes.flatten()
for i, col in enumerate(numerical_cols):
    sns.regplot(ax=axes[i], data=valve_data, x=col, y='Lifespan',
               scatter_kws={'s': 10}, line_kws={'color': 'red'})
    axes[i].set_title(f'Valve: Lifespan vs {col}')
    axes[i].set_xlabel(col)
    axes[i].set_ylabel('Lifespan (hours)')
for j in range(len(numerical_cols), len(axes)):
    fig.delaxes(axes[j])
plt.tight_layout()
plt.show()
```

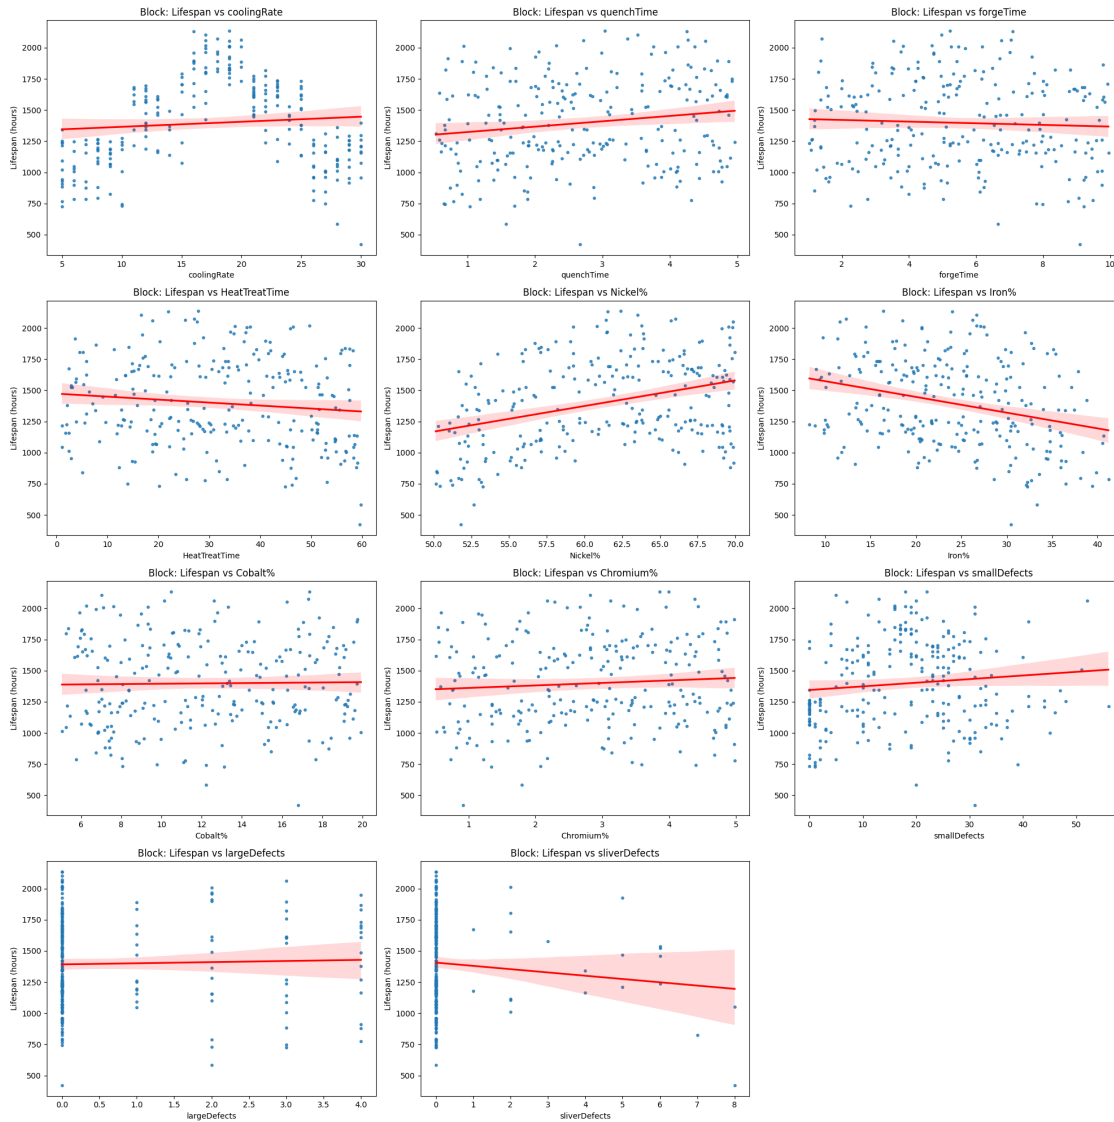


```
[39]: # 1. Box Plots for Valve
fig, axes = plt.subplots(nrows=2, ncols=2, figsize=(16, 12))
axes = axes.flatten()
for i, col in enumerate(['microstructure', 'seedLocation', 'castType']):
    sns.boxplot(ax=axes[i], x=col, y='Lifespan', data=valve_data)
    axes[i].set_title(f'Valve: Lifespan by {col}')
    axes[i].set_xlabel(col)
    axes[i].set_ylabel('Lifespan (hours)')
    axes[i].tick_params(axis='x', rotation=45)
for j in range(len(['microstructure', 'seedLocation', 'castType']), len(axes)):
    fig.delaxes(axes[j])
plt.tight_layout()
plt.show()
```

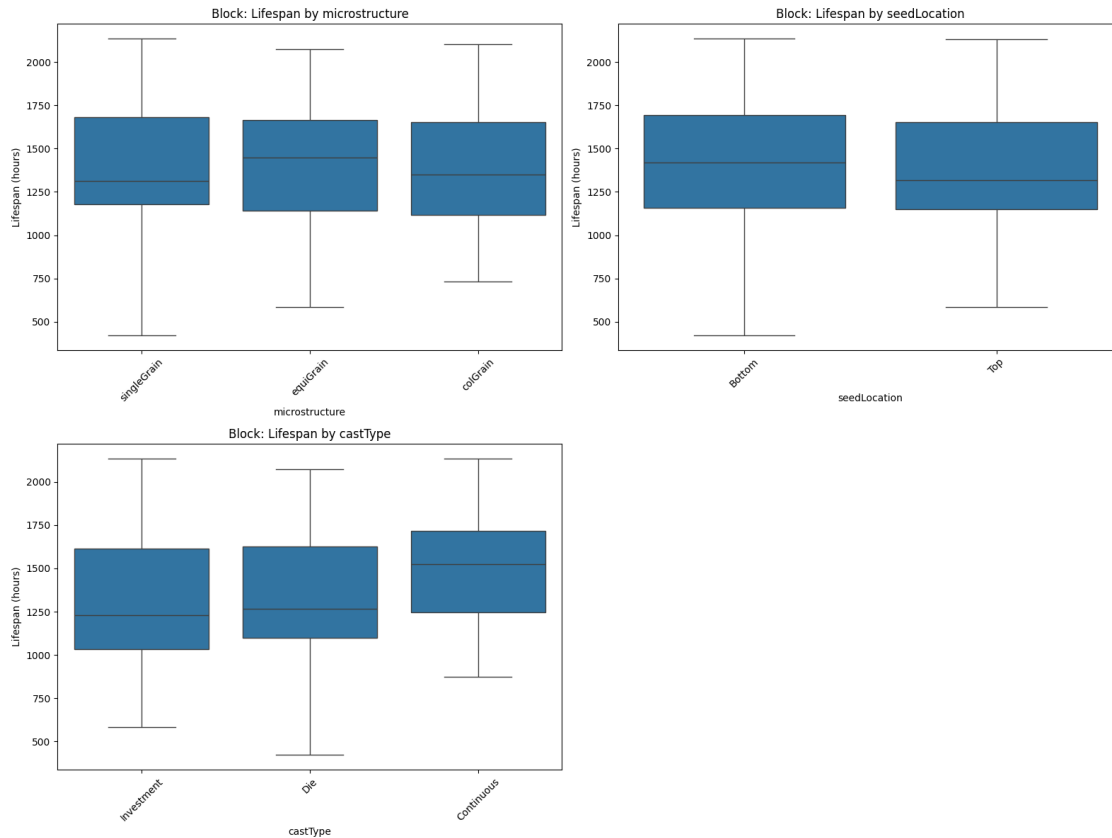


1.4.2 Block

```
[35]: # 2. Scatter Plots for Block
block_data = df[df['partType'] == 'Block']
fig, axes = plt.subplots(nrows=4, ncols=3, figsize=(20, 20))
axes = axes.flatten()
for i, col in enumerate(numerical_cols):
    sns.regplot(ax=axes[i], data=block_data, x=col, y='Lifespan',
               scatter_kws={'s': 10}, line_kws={'color': 'red'})
    axes[i].set_title(f'Block: Lifespan vs {col}')
    axes[i].set_xlabel(col)
    axes[i].set_ylabel('Lifespan (hours)')
for j in range(len(numerical_cols), len(axes)):
    fig.delaxes(axes[j])
plt.tight_layout()
plt.show()
```

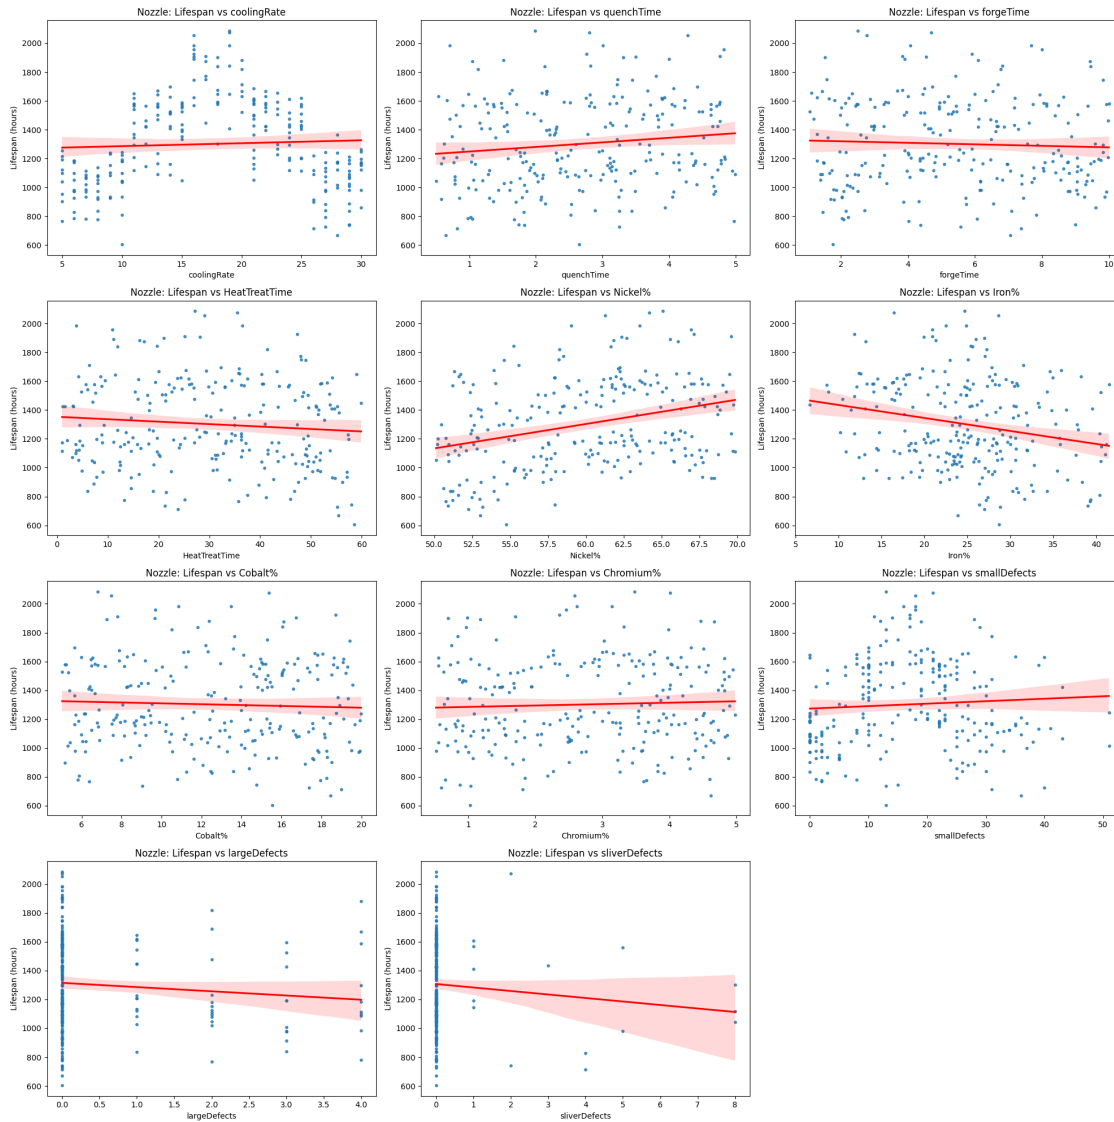


```
[40]: # 2. Box Plots for Block
fig, axes = plt.subplots(nrows=2, ncols=2, figsize=(16, 12))
axes = axes.flatten()
for i, col in enumerate(['microstructure', 'seedLocation', 'castType']):
    sns.boxplot(ax=axes[i], x=col, y='Lifespan', data=block_data)
    axes[i].set_title(f'Block: Lifespan by {col}')
    axes[i].set_xlabel(col)
    axes[i].set_ylabel('Lifespan (hours)')
    axes[i].tick_params(axis='x', rotation=45)
for j in range(len(['microstructure', 'seedLocation', 'castType']), len(axes)):
    fig.delaxes(axes[j])
plt.tight_layout()
plt.show()
```

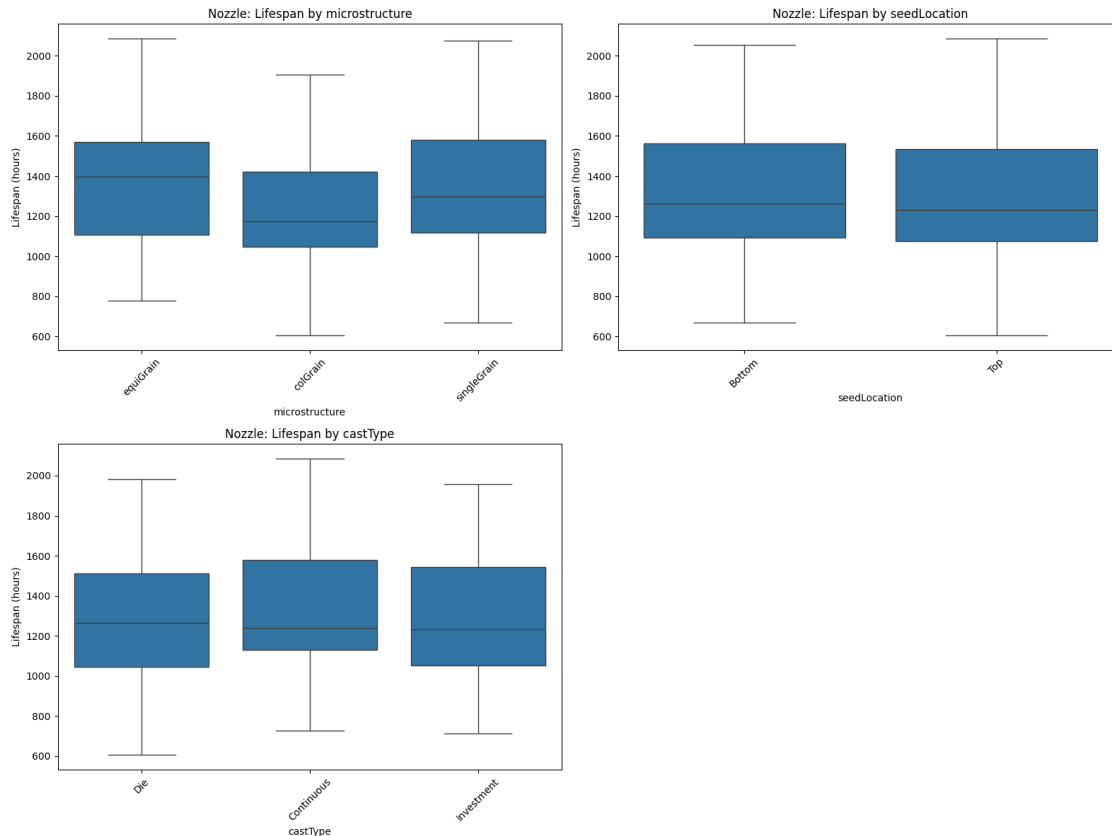


1.4.3 Nozzle

```
[36]: # 3. Scatter Plots for Nozzle
nozzle_data = df[df['partType'] == 'Nozzle']
fig, axes = plt.subplots(nrows=4, ncols=3, figsize=(20, 20))
axes = axes.flatten()
for i, col in enumerate(numerical_cols):
    sns.regplot(ax=axes[i], data=nozzle_data, x=col, y='Lifespan',
               scatter_kws={'s': 10}, line_kws={'color': 'red'})
    axes[i].set_title(f'Nozzle: Lifespan vs {col}')
    axes[i].set_xlabel(col)
    axes[i].set_ylabel('Lifespan (hours)')
for j in range(len(numerical_cols), len(axes)):
    fig.delaxes(axes[j])
plt.tight_layout()
plt.show()
```

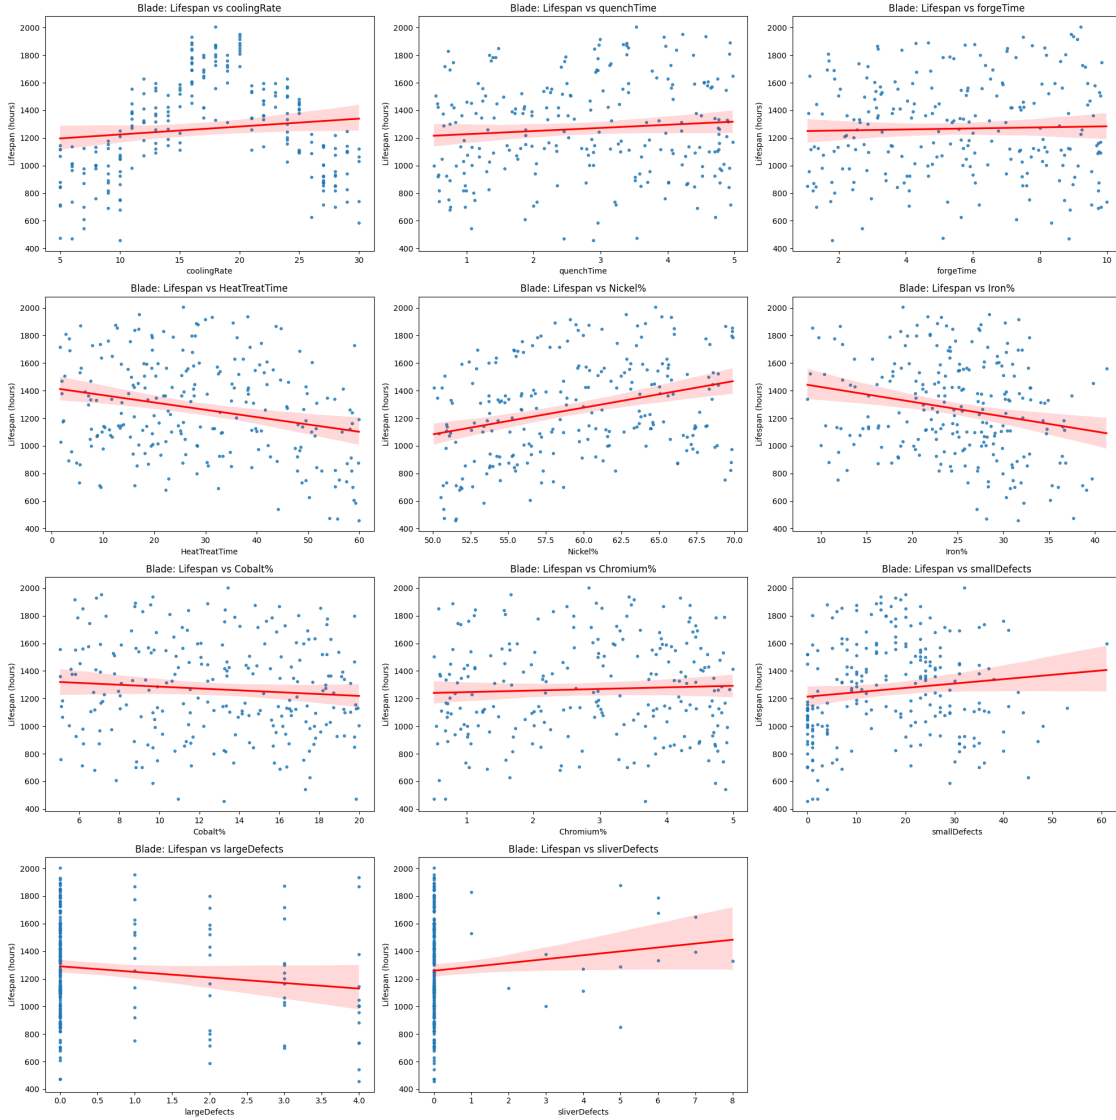


```
[41]: # 3. Box Plots for Nozzle
fig, axes = plt.subplots(nrows=2, ncols=2, figsize=(16, 12))
axes = axes.flatten()
for i, col in enumerate(['microstructure', 'seedLocation', 'castType']):
    sns.boxplot(ax=axes[i], x=col, y='Lifespan', data=nozzle_data)
    axes[i].set_title(f'Nozzle: Lifespan by {col}')
    axes[i].set_xlabel(col)
    axes[i].set_ylabel('Lifespan (hours)')
    axes[i].tick_params(axis='x', rotation=45)
for j in range(len(['microstructure', 'seedLocation', 'castType']), len(axes)):
    fig.delaxes(axes[j])
plt.tight_layout()
plt.show()
```

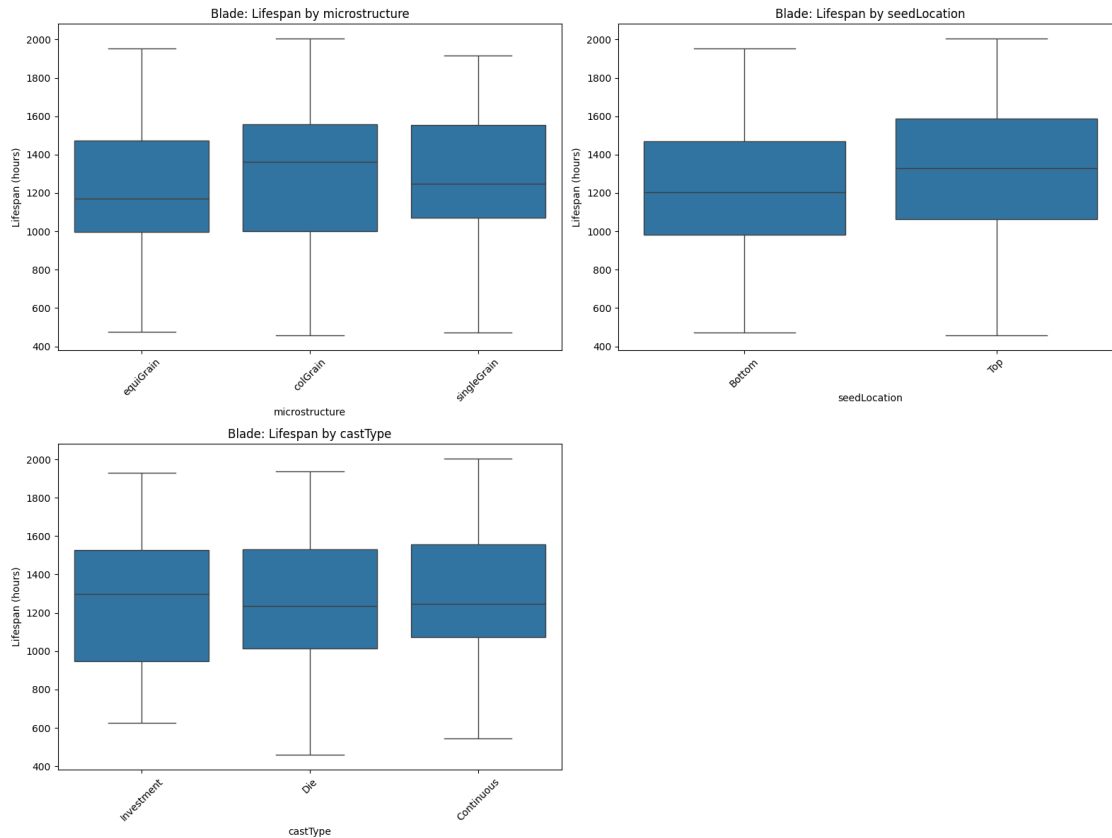


1.4.4 Blade

```
[37]: # 4. Scatter Plots for Blade
blade_data = df[df['partType'] == 'Blade']
fig, axes = plt.subplots(nrows=4, ncols=3, figsize=(20, 20))
axes = axes.flatten()
for i, col in enumerate(numerical_cols):
    sns.regplot(ax=axes[i], data=blade_data, x=col, y='Lifespan',
                scatter_kws={'s': 10}, line_kws={'color': 'red'})
    axes[i].set_title(f'Blade: Lifespan vs {col}')
    axes[i].set_xlabel(col)
    axes[i].set_ylabel('Lifespan (hours)')
for j in range(len(numerical_cols), len(axes)):
    fig.delaxes(axes[j])
plt.tight_layout()
plt.show()
```




```
[42]: # 4. Box Plots for Blade
fig, axes = plt.subplots(nrows=2, ncols=2, figsize=(16, 12))
axes = axes.flatten()
for i, col in enumerate(['microstructure', 'seedLocation', 'castType']):
    sns.boxplot(ax=axes[i], x=col, y='Lifespan', data=blade_data)
    axes[i].set_title(f'Blade: Lifespan by {col}')
    axes[i].set_xlabel(col)
    axes[i].set_ylabel('Lifespan (hours)')
    axes[i].tick_params(axis='x', rotation=45)
for j in range(len(['microstructure', 'seedLocation', 'castType']), len(axes)):
    fig.delaxes(axes[j])
plt.tight_layout()
plt.show()
```

1.5 Correlation Matrix for Combined Lifespan of Metal Parts

```
[43]: # Include 'Lifespan' in the correlation matrix
numerical_cols_with_target = numerical_cols + ['Lifespan']
corr_matrix = df[numerical_cols_with_target].corr()

# Plot heatmap
plt.figure(figsize=(12, 8))
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm')
plt.title('Correlation Matrix')
plt.show()
```



1.6 Correlation Matrix for Lifespan of Individual Metal Parts

```
[44]: # Correlation Matrices for All Metal Parts in Subplots
fig, axes = plt.subplots(nrows=2, ncols=2, figsize=(24, 16))
axes = axes.flatten()

# 1. Correlation Matrix for Valve
valve_corr_matrix = valve_data[numerical_cols + ['Lifespan']].corr()
sns.heatmap(valve_corr_matrix, annot=True, cmap='coolwarm', ax=axes[0])
axes[0].set_title('Correlation Matrix for Valve')

# 2. Correlation Matrix for Block
block_corr_matrix = block_data[numerical_cols + ['Lifespan']].corr()
sns.heatmap(block_corr_matrix, annot=True, cmap='coolwarm', ax=axes[1])
axes[1].set_title('Correlation Matrix for Block')

# 3. Correlation Matrix for Nozzle
nozzle_corr_matrix = nozzle_data[numerical_cols + ['Lifespan']].corr()
sns.heatmap(nozzle_corr_matrix, annot=True, cmap='coolwarm', ax=axes[2])
axes[2].set_title('Correlation Matrix for Nozzle')
```

4. Correlation Matrix for Blade

```
blade_corr_matrix = blade_data[numerical_cols + ['Lifespan']].corr()
sns.heatmap(blade_corr_matrix, annot=True, cmap='coolwarm', ax=axes[3])
axes[3].set_title('Correlation Matrix for Blade')
```

```
plt.tight_layout()
plt.show()
```

