

Testing

November 7, 2024

1 COMP1801 - Machine Learning Coursework Solution

Let's start by importing the essential Python libraries for data analysis and machine learning.

```
[151]: # Import libraries
try:
    # Importing general libraries
    import os
    import glob
    import pandas as pd
    import joblib

    # Importing libraries for data visualization
    import matplotlib.pyplot as plt
    import numpy as np

    # Importing libraries for model building
    from sklearn.preprocessing import LabelEncoder
    from sklearn.model_selection import train_test_split, GridSearchCV,
↳RandomizedSearchCV
    from sklearn.ensemble import RandomForestRegressor
    from sklearn.metrics import root_mean_squared_error, r2_score,
↳mean_absolute_error, mean_squared_log_error

    # Importing libraries for data preprocessing
    from scipy.stats import randint

except Exception as e:
    print(f"Error : {e}")
```

```
[152]: # Find the CSV file in the Datasets directory
data_path = '../Datasets/*.csv'
file_list = glob.glob(data_path)

for file in file_list:
    print(f"Found file: {file}")

# Ensure there is exactly one file
```

```

if len(file_list) == 1:
    # Load the dataset
    df = pd.read_csv(file_list[0])
    print(f"Loaded dataset: {file_list[0]}")
else:
    raise FileNotFoundError("No CSV file found or multiple CSV files found in_
↳the Datasets directory.")

```

Found file: ../Datasets/Dataset.csv

Loaded dataset: ../Datasets/Dataset.csv

```

[153]: # File path to save the trained model
destination = '../Models/'
os.makedirs(destination, exist_ok=True)

```

```

[154]: df.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 16 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   Lifespan              1000 non-null   float64
 1   partType              1000 non-null   object
 2   microstructure        1000 non-null   object
 3   coolingRate            1000 non-null   int64
 4   quenchTime            1000 non-null   float64
 5   forgeTime             1000 non-null   float64
 6   HeatTreatTime         1000 non-null   float64
 7   Nickel%               1000 non-null   float64
 8   Iron%                 1000 non-null   float64
 9   Cobalt%               1000 non-null   float64
10   Chromium%             1000 non-null   float64
11   smallDefects          1000 non-null   int64
12   largeDefects          1000 non-null   int64
13   sliverDefects         1000 non-null   int64
14   seedLocation          1000 non-null   object
15   castType              1000 non-null   object
dtypes: float64(8), int64(4), object(4)
memory usage: 125.1+ KB

```

```

[155]: # Check for missing values
df.isnull().sum()

```

```

[155]: Lifespan          0
partType              0
microstructure        0
coolingRate           0

```

```

quenchTime      0
forgeTime       0
HeatTreatTime   0
Nickel%         0
Iron%           0
Cobalt%         0
Chromium%       0
smallDefects    0
largeDefects    0
sliverDefects   0
seedLocation    0
castType        0
dtype: int64

```

```
[156]: df.head()
```

```

[156]:   Lifespan  partType  microstructure  coolingRate  quenchTime  forgeTime  \
0    1469.17   Nozzle    equiGrain         13         3.84         6.47
1    1793.64   Block    singleGrain        19         2.62         3.48
2     700.60   Blade    equiGrain         28         0.76         1.34
3    1082.10   Nozzle    colGrain          9         2.01         2.19
4    1838.83   Blade    colGrain         16         4.13         3.87

      HeatTreatTime  Nickel%  Iron%  Cobalt%  Chromium%  smallDefects  \
0          46.87     65.73  16.52   16.82         0.93           10
1          44.70     54.22  35.38    6.14         4.26           19
2           9.54     51.83  35.95    8.81         3.41           35
3          20.29     57.03  23.33   16.86         2.78            0
4          16.13     59.62  27.37   11.45         1.56           10

      largeDefects  sliverDefects  seedLocation  castType
0                0              0      Bottom      Die
1                0              0      Bottom  Investment
2                3              0      Bottom  Investment
3                1              0          Top  Continuous
4                0              0          Top      Die

```

```
[157]: df.describe()
```

```

[157]:   Lifespan  coolingRate  quenchTime  forgeTime  HeatTreatTime  \
count  1000.000000  1000.000000  1000.000000  1000.000000  1000.000000
mean    1298.556320    17.639000    2.764230    5.464600    30.194510
std     340.071434     7.491783    1.316979    2.604513   16.889415
min     417.990000     5.000000    0.500000    1.030000     1.030000
25%    1047.257500    11.000000    1.640000    3.170000   16.185000
50%    1266.040000    18.000000    2.755000    5.475000   29.365000
75%    1563.050000    24.000000    3.970000    7.740000   44.955000

```

max	2134.530000	30.000000	4.990000	10.000000	59.910000
-----	-------------	-----------	----------	-----------	-----------

	Nickel%	Iron%	Cobalt%	Chromium%	smallDefects \
count	1000.000000	1000.000000	1000.000000	1000.000000	1000.000000
mean	60.243080	24.553580	12.434690	2.768650	17.311000
std	5.790475	7.371737	4.333197	1.326496	12.268365
min	50.020000	6.660000	5.020000	0.510000	0.000000
25%	55.287500	19.387500	8.597500	1.590000	7.000000
50%	60.615000	24.690000	12.585000	2.865000	18.000000
75%	65.220000	29.882500	16.080000	3.922500	26.000000
max	69.950000	43.650000	19.990000	4.990000	61.000000

	largeDefects	sliverDefects
count	1000.000000	1000.000000
mean	0.550000	0.292000
std	1.163982	1.199239
min	0.000000	0.000000
25%	0.000000	0.000000
50%	0.000000	0.000000
75%	0.000000	0.000000
max	4.000000	8.000000

```
[158]: # Using nunique()
num_parts = df['partType'].nunique()
print(f"Number of unique parts types: {num_parts}")

# Or using value_counts() to see the distribution
parts_distribution = df['partType'].value_counts()
print("\nDistribution of parts types:")
print(parts_distribution)
```

Number of unique parts types: 4

Distribution of parts types:

```
partType
Valve      265
Block      253
Nozzle     245
Blade      237
Name: count, dtype: int64
```

```
[159]: categorical_cols_unified = ['partType', 'microstructure', 'seedLocation',
    ↪ 'castType']

# Create a DataFrame to display unique values and their counts
unique_values_df = pd.DataFrame({
    'Column': categorical_cols_unified,
```

```

    'Unique Values': [df[col].unique().tolist() for col in
↪categorical_cols_unified],
    'Count of Unique Values': [df[col].nunique() for col in
↪categorical_cols_unified]
})

print(unique_values_df)

```

	Column	Unique Values	Count of Unique Values
0	partType	[Nozzle, Block, Blade, Valve]	4
1	microstructure	[equiGrain, singleGrain, colGrain]	3
2	seedLocation	[Bottom, Top]	2
3	castType	[Die, Investment, Continuous]	3

```

[160]: # Creating a copy of the dataframe to ensure we maintain the original intact
df_label_encoded = df.copy()
encoder = "Label Encoding"

# Apply Label Encoding to each categorical column
label_encoders = {}
for col in categorical_cols_unified:
    le = LabelEncoder()
    df_label_encoded[col] = le.fit_transform(df_label_encoded[col])
    label_encoders[col] = le # Store the encoder for inverse transformation if
↪needed later

# Display the first few rows to verify
display(df_label_encoded.head())

```

	Lifespan	partType	microstructure	coolingRate	quenchTime	forgeTime	\
0	1469.17	2	1	13	3.84	6.47	
1	1793.64	1	2	19	2.62	3.48	
2	700.60	0	1	28	0.76	1.34	
3	1082.10	2	0	9	2.01	2.19	
4	1838.83	0	0	16	4.13	3.87	

	HeatTreatTime	Nickel%	Iron%	Cobalt%	Chromium%	smallDefects	\
0	46.87	65.73	16.52	16.82	0.93	10	
1	44.70	54.22	35.38	6.14	4.26	19	
2	9.54	51.83	35.95	8.81	3.41	35	
3	20.29	57.03	23.33	16.86	2.78	0	
4	16.13	59.62	27.37	11.45	1.56	10	

	largeDefects	sliverDefects	seedLocation	castType
0	0	0	0	1
1	0	0	0	2
2	3	0	0	2
3	1	0	1	0

4 0 0 1 1

[161]: df_label_encoded.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 16 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Lifespan              1000 non-null   float64
1   partType              1000 non-null   int64
2   microstructure        1000 non-null   int64
3   coolingRate            1000 non-null   int64
4   quenchTime            1000 non-null   float64
5   forgeTime             1000 non-null   float64
6   HeatTreatTime         1000 non-null   float64
7   Nickel%               1000 non-null   float64
8   Iron%                 1000 non-null   float64
9   Cobalt%               1000 non-null   float64
10  Chromium%             1000 non-null   float64
11  smallDefects           1000 non-null   int64
12  largeDefects           1000 non-null   int64
13  sliverDefects          1000 non-null   int64
14  seedLocation           1000 non-null   int64
15  castType               1000 non-null   int64
dtypes: float64(8), int64(8)
memory usage: 125.1 KB
```

```
[162]: # Define the target variable and feature set
X = df_label_encoded.drop(columns=['Lifespan']) # Features
y = df_label_encoded['Lifespan'] # Target

# Split the dataset into training and testing sets (80% train, 20% test)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳ random_state=42)

# Display the shapes of the training and testing sets to verify
print(f'--- {encoder} Shape ---\n')
print("X_train shape:", X_train.shape)
print("X_test shape:", X_test.shape)
print("y_train shape:", y_train.shape)
print("y_test shape:", y_test.shape)
```

--- Label Encoding Shape ---

```
X_train shape: (800, 15)
X_test shape: (200, 15)
y_train shape: (800,)
y_test shape: (200,)
```

```
[163]: # Initialize the Random Forest Regressor with default parameters
rf_model = RandomForestRegressor(n_estimators=100, random_state=42)

# Fit the model to the training data
rf_model.fit(X_train, y_train)

# Make predictions on the test set
y_pred = rf_model.predict(X_test)

# Evaluate the model using RMSE, R2 Score, and MAE
rmse = root_mean_squared_error(y_test, y_pred) # Root Mean Squared Error
r2 = r2_score(y_test, y_pred) # R2 Score
mae = mean_absolute_error(y_test, y_pred) # Mean Absolute Error
msle = mean_squared_log_error(y_test, y_pred) # Mean Squared Log Error

print(f'--- {encoder} Performance ---\n')
print(f"Root Mean Squared Error (RMSE): {rmse:.2f}")
print(f"R2 Score: {r2:.2f}")
print(f"Mean Absolute Error (MAE): {mae:.2f}")
print(f"Mean Squared Log Error (MSLE): {msle:.2f}")
```

--- Label Encoding Performance ---

Root Mean Squared Error (RMSE): 90.95
R² Score: 0.92
Mean Absolute Error (MAE): 72.50
Mean Squared Log Error (MSLE): 0.01

```
[164]: print(f"\nFeatures saved with {encoder}\n")
for col in X.columns:
    print(f"- {col}")
```

Features saved with Label Encoding

- partType
- microstructure
- coolingRate
- quenchTime
- forgeTime
- HeatTreatTime
- Nickel%
- Iron%
- Cobalt%
- Chromium%
- smallDefects
- largeDefects
- sliverDefects

- seedLocation
- castType

```
[165]: # Features to drop based on low importance for Label Encoding
low_importance_features = ['seedLocation', 'microstructure', 'castType',
    ↪ 'smallDefects', 'sliverDefects']
```

```
[166]: print(f"--- {encoder} Shape Reduction ---")

# Split the dataset into training and testing sets (80% train, 20% test)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    ↪ random_state=42)

# Display the shapes of the training and testing sets to verify
print("\nOriginal Shapes")
print("X_train shape:", X_train.shape)
print("X_test shape:", X_test.shape)
print("y_train shape:", y_train.shape)
print("y_test shape:", y_test.shape)

# Features to drop based on low importance
low_importance_features = ['seedLocation', 'microstructure', 'castType',
    ↪ 'smallDefects', 'sliverDefects']

# Create a new DataFrame excluding these features
X_train_reduced = X_train.drop(columns=low_importance_features)
X_test_reduced = X_test.drop(columns=low_importance_features)

# Display the shapes of the training and testing sets to verify
print("\nUpdated Shapes")
print("X_train_reduced shape:", X_train_reduced.shape)
print("X_test_reduced shape:", X_test_reduced.shape)
print("y_train shape:", y_train.shape)
print("y_test shape:", y_test.shape)
```

--- Label Encoding Shape Reduction ---

Original Shapes

X_train shape: (800, 15)
X_test shape: (200, 15)
y_train shape: (800,)
y_test shape: (200,)

Updated Shapes

X_train_reduced shape: (800, 10)
X_test_reduced shape: (200, 10)
y_train shape: (800,)
y_test shape: (200,)


```
[167]: # Initialize the Random Forest Regressor with best parameters
rf_model = RandomForestRegressor(
    max_depth=15,
    n_estimators=387,
    random_state=42
)

# Fit the model to the training data
rf_model.fit(X_train_reduced, y_train)

# Make predictions on the test set
y_pred = rf_model.predict(X_test_reduced)

# Evaluate the model using RMSE, R2 Score, and MAE
rmse = root_mean_squared_error(y_test, y_pred) # Root Mean Squared Error
r2 = r2_score(y_test, y_pred) # R2 Score
mae = mean_absolute_error(y_test, y_pred) # Mean Absolute Error
msle = mean_squared_log_error(y_test, y_pred) # Mean Squared Log Error

print(f"Root Mean Squared Error (RMSE): {rmse:.2f}")
print(f"R2 Score: {r2:.2f}")
print(f"Mean Absolute Error (MAE): {mae:.2f}")
print(f"Mean Squared Log Error (MSLE): {msle:.2f}")
```

Root Mean Squared Error (RMSE): 88.10
R² Score: 0.93
Mean Absolute Error (MAE): 70.68
Mean Squared Log Error (MSLE): 0.01

```
[168]: # Creating a copy of the dataframe to ensure we maintain the original intact
df_onehot_encoded = df.copy()
encoder = "One-Hot Encoding"

# Apply one-hot encoding to the categorical columns
df_onehot_encoded = pd.get_dummies(df_onehot_encoded,
    columns=categorical_cols_unified, drop_first=False)

# Display the first few rows to verify
display(df_onehot_encoded.head())
```

	Lifespan	coolingRate	quenchTime	forgeTime	HeatTreatTime	Nickel%	\
0	1469.17	13	3.84	6.47	46.87	65.73	
1	1793.64	19	2.62	3.48	44.70	54.22	
2	700.60	28	0.76	1.34	9.54	51.83	
3	1082.10	9	2.01	2.19	20.29	57.03	
4	1838.83	16	4.13	3.87	16.13	59.62	

	Iron%	Cobalt%	Chromium%	smallDefects	...	partType_Nozzle	\
--	-------	---------	-----------	--------------	-----	-----------------	---

0	16.52	16.82	0.93	10	...	True
1	35.38	6.14	4.26	19	...	False
2	35.95	8.81	3.41	35	...	False
3	23.33	16.86	2.78	0	...	True
4	27.37	11.45	1.56	10	...	False

	partType_Valve	microstructure_colGrain	microstructure_equiGrain	\
0	False	False	True	
1	False	False	False	
2	False	False	True	
3	False	True	False	
4	False	True	False	

	microstructure_singleGrain	seedLocation_Bottom	seedLocation_Top	\
0	False	True	False	
1	True	True	False	
2	False	True	False	
3	False	False	True	
4	False	False	True	

	castType_Continuous	castType_Die	castType_Investment
0	False	True	False
1	False	False	True
2	False	False	True
3	True	False	False
4	False	True	False

[5 rows x 24 columns]

```
[169]: df_onehot_encoded.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 1000 entries, 0 to 999
```

```
Data columns (total 24 columns):
```

#	Column	Non-Null Count	Dtype
---	-----	-----	-----
0	Lifespan	1000 non-null	float64
1	coolingRate	1000 non-null	int64
2	quenchTime	1000 non-null	float64
3	forgeTime	1000 non-null	float64
4	HeatTreatTime	1000 non-null	float64
5	Nickel%	1000 non-null	float64
6	Iron%	1000 non-null	float64
7	Cobalt%	1000 non-null	float64
8	Chromium%	1000 non-null	float64
9	smallDefects	1000 non-null	int64
10	largeDefects	1000 non-null	int64
11	sliverDefects	1000 non-null	int64

```

12 partType_Blade          1000 non-null    bool
13 partType_Block          1000 non-null    bool
14 partType_Nozzle         1000 non-null    bool
15 partType_Valve          1000 non-null    bool
16 microstructure_colGrain  1000 non-null    bool
17 microstructure_equiGrain 1000 non-null    bool
18 microstructure_singleGrain 1000 non-null    bool
19 seedLocation_Bottom     1000 non-null    bool
20 seedLocation_Top        1000 non-null    bool
21 castType_Continuous     1000 non-null    bool
22 castType_Die            1000 non-null    bool
23 castType_Investment     1000 non-null    bool
dtypes: bool(12), float64(8), int64(4)
memory usage: 105.6 KB

```

```

[170]: # Define the target variable and feature set
X = df_onehot_encoded.drop(columns=['Lifespan']) # Features
y = df_onehot_encoded['Lifespan'] # Target

# Split the dataset into training and testing sets (80% train, 20% test)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳ random_state=42)

# Display the shapes of the training and testing sets to verify
print(f'--- {encoder} Shape ---\n')
print("X_train shape:", X_train.shape)
print("X_test shape:", X_test.shape)
print("y_train shape:", y_train.shape)
print("y_test shape:", y_test.shape)

```

--- One-Hot Encoding Shape ---

```

X_train shape: (800, 23)
X_test shape: (200, 23)
y_train shape: (800,)
y_test shape: (200,)

```

```

[171]: # Initialize the Random Forest Regressor with default parameters
rf_model = RandomForestRegressor(n_estimators=100, random_state=42)

# Fit the model to the training data
rf_model.fit(X_train, y_train)

# Make predictions on the test set
y_pred = rf_model.predict(X_test)

# Evaluate the model using RMSE, R² Score, and MAE
rmse = root_mean_squared_error(y_test, y_pred) # Root Mean Squared Error

```

```

r2 = r2_score(y_test, y_pred) # R2 Score
mae = mean_absolute_error(y_test, y_pred) # Mean Absolute Error
msle = mean_squared_log_error(y_test, y_pred) # Mean Squared Log Error

print(f'--- {encoder} Performance ---\n')
print(f"Root Mean Squared Error (RMSE): {rmse:.2f}")
print(f"R2 Score: {r2:.2f}")
print(f"Mean Absolute Error (MAE): {mae:.2f}")
print(f"Mean Squared Log Error (MSLE): {msle:.2f}")

```

--- One-Hot Encoding Performance ---

Root Mean Squared Error (RMSE): 85.15
 R^2 Score: 0.93
Mean Absolute Error (MAE): 67.46
Mean Squared Log Error (MSLE): 0.01

```

[174]: # Train the final model with the best parameters with the best parameters
final_rf_model = RandomForestRegressor(
    max_depth=15,
    n_estimators=387,
    random_state=42
)

print(f'--- Generating model with {encoder} ---')

# Fit the model to the training data
final_rf_model.fit(X_train, y_train)

# Save the trained model
joblib.dump(final_rf_model, f'{destination}Final-RFModel.pkl')

```

--- Generating model with One-Hot Encoding ---

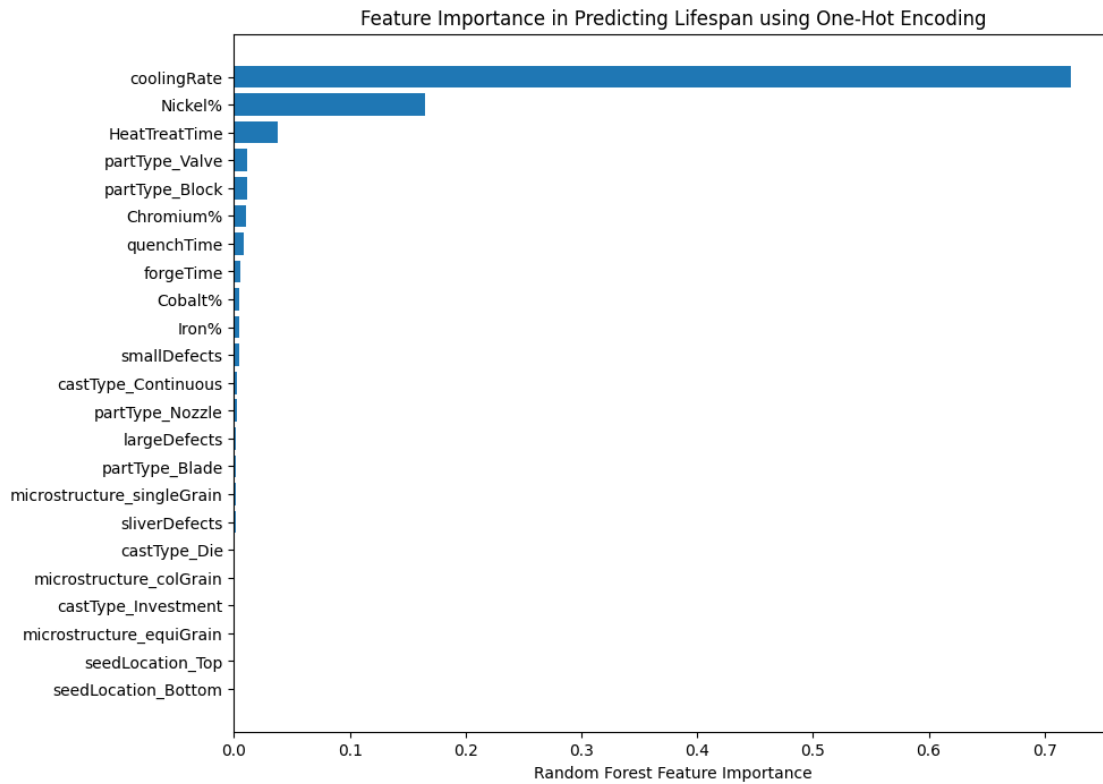
```
[174]: ['../Models/Final-RFModel.pkl']
```

```

[175]: # Feature importance analysis
feature_importances = final_rf_model.feature_importances_
feature_names = X_train.columns

# Sort features by importance
sorted_idx = np.argsort(feature_importances)
plt.figure(figsize=(10, 8))
plt.barh(feature_names[sorted_idx], feature_importances[sorted_idx])
plt.xlabel("Random Forest Feature Importance")
plt.title(f"Feature Importance in Predicting Lifespan using {encoder}")
plt.show()

```



```
[176]: print(f"\nFeatures saved with {encoder}\n")
      for col in X.columns:
          print(f"- {col}")
```

Features saved with One-Hot Encoding

- coolingRate
- quenchTime
- forgeTime
- HeatTreatTime
- Nickel%
- Iron%
- Cobalt%
- Chromium%
- smallDefects
- largeDefects
- sliverDefects
- partType_Blade
- partType_Block
- partType_Nozzle
- partType_Valve

- microstructure_colGrain
- microstructure_equiGrain
- microstructure_singleGrain
- seedLocation_Bottom
- seedLocation_Top
- castType_Continuous
- castType_Die
- castType_Investment

```
[177]: # Selecting only the most important features (those with non-zero or
↳significant importance)
important_features = ['partType_Blade', 'partType_Block', 'partType_Nozzle',
↳'partType_Valve', 'coolingRate', 'Nickel%', 'HeatTreatTime', 'Chromium%',
↳'quenchTime']
X_important = df_onehot_encoded[important_features]

print(f'--- {encoder} Updated Performance ---\n')

# Split the dataset and retrain the model
X_train_imp, X_test_imp, y_train_imp, y_test_imp =
↳train_test_split(X_important, y, test_size=0.2, random_state=42)

# Initialize and train the Random Forest Regressor with reduced features with
↳the best parameters
rf_model_imp = RandomForestRegressor(
    max_depth=15,
    n_estimators=387,
    random_state=42
)
rf_model_imp.fit(X_train_imp, y_train_imp)

# Evaluate the model
y_pred_imp = rf_model_imp.predict(X_test_imp)
rmse_imp = root_mean_squared_error(y_test_imp, y_pred_imp)
r2_imp = r2_score(y_test_imp, y_pred_imp)
mae_imp = mean_absolute_error(y_test_imp, y_pred_imp)
msle_imp = mean_squared_log_error(y_test_imp, y_pred_imp)

print(f"Reduced Features RMSE: {rmse_imp:.2f}")
print(f"Reduced Features R2 Score: {r2_imp:.2f}")
print(f"Reduced Features MAE: {mae_imp:.2f}")
print(f"Reduced Features MSLE: {msle_imp:.2f}")
```

--- One-Hot Encoding Updated Performance ---

Reduced Features RMSE: 78.70
Reduced Features R² Score: 0.94
Reduced Features MAE: 61.14


Reduced Features MSLE: 0.01

```
[178]: import sklearn
print(sklearn.__version__)
```

1.5.2

```
[179]: from sklearn.inspection import PartialDependenceDisplay

# Plot partial dependence for important features using the new API
features_to_plot = ['coolingRate', 'Nickel%']
PartialDependenceDisplay.from_estimator(rf_model_imp, X_train_imp,
    features_to_plot, grid_resolution=50)
plt.show()
```



```
[180]: from sklearn.model_selection import cross_val_score
from sklearn.metrics import make_scorer, root_mean_squared_error

# Define the model with best-found hyperparameters
rf_model_cv = RandomForestRegressor(
    max_depth=15,
    n_estimators=387,
```

```

        min_samples_leaf=2,
        min_samples_split=3,
        random_state=42
    )

    # Define the scoring metric: RMSE
    scorer = make_scorer(root_mean_squared_error)

    # Perform 5-Fold Cross-Validation
    cv_scores = cross_val_score(rf_model_cv, X_important, y, cv=5, scoring=scorer)

    # Print results
    print("Cross-Validation RMSE scores for each fold: ", cv_scores)
    print(f"Mean CV RMSE: {cv_scores.mean():.2f}")
    print(f"Standard Deviation of CV RMSE: {cv_scores.std():.2f}")

```

Cross-Validation RMSE scores for each fold: [77.55548306 74.97658535
77.56023918 82.12528196 81.38495862]
Mean CV RMSE: 78.72
Standard Deviation of CV RMSE: 2.66

```

[181]: from sklearn.compose import ColumnTransformer
        from sklearn.preprocessing import OneHotEncoder, LabelEncoder
        from sklearn.pipeline import Pipeline
        from sklearn.ensemble import RandomForestRegressor

        # Separate features for encoding
        onehot_features = ['microstructure', 'seedLocation', 'castType']
        label_features = ['partType']

        # Custom transformers
        label_encoder = LabelEncoder()
        df['partType'] = label_encoder.fit_transform(df['partType']) # Apply label_
                               ↪ encoding directly

        # Create a ColumnTransformer for One-Hot Encoding
        preprocessor = ColumnTransformer(
            transformers=[
                # To match the output of One-Hot Encoding with the Label Encoding
                ('onehot', OneHotEncoder(dtype=int), onehot_features)
            ],
            remainder='passthrough' # Keeps all other features as they are
        )

        # Define pipeline with preprocessing and model
        pipeline = Pipeline([
            ('preprocessor', preprocessor),

```



```

    ('model', RandomForestRegressor(
        max_depth=15,
        n_estimators=387,
        min_samples_leaf=2,
        min_samples_split=3,
        random_state=42
    ))
])

# Split dataset
X_train, X_test, y_train, y_test = train_test_split(df.
    ↳drop(columns=['Lifespan']), df['Lifespan'], test_size=0.2, random_state=42)

# Fit the pipeline
pipeline.fit(X_train, y_train)

# Evaluate the model
y_pred = pipeline.predict(X_test)
rmse = root_mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)
mae = mean_absolute_error(y_test, y_pred)
msle = mean_squared_log_error(y_test, y_pred)

print(f"RMSE: {rmse:.2f}")
print(f"R2 Score: {r2:.2f}")
print(f"MAE: {mae:.2f}")
print(f"MSLE: {msle:.2f}")

```

RMSE: 89.91
 R² Score: 0.92
 MAE: 71.64
 MSLE: 0.01

```

[182]: import matplotlib.pyplot as plt
import numpy as np

# Step 1: Access the trained Random Forest model from the pipeline
rf_model_imp = pipeline.named_steps['model']

# Step 2: Retrieve the feature names from the preprocessor
# Get the One-Hot Encoder categories and append label encoded features
onehot_encoder = pipeline.named_steps['preprocessor'].
    ↳named_transformers_['onehot']
onehot_feature_names = onehot_encoder.get_feature_names_out(onehot_features)

# Combine One-Hot Encoded feature names with label-encoded and numerical
    ↳feature names

```

```


all_feature_names = list(onehot_feature_names) + label_features + [col for col
    ↪ in X_train.columns if col not in onehot_features + label_features]

# Step 3: Extract the feature importances from the RandomForestRegressor
importances = rf_model_imp.feature_importances_

# Sort feature importances in descending order
indices = np.argsort(importances)[::-1]

# Step 4: Plot the feature importances
plt.figure(figsize=(10, 8))
plt.title("Feature Importance after Hybrid Encoding")
plt.bar(range(len(importances)), importances[indices], align="center")
plt.xticks(range(len(importances)), np.array(all_feature_names)[indices],
    ↪ rotation=90)
plt.xlabel("Features")
plt.ylabel("Importance Score")
plt.tight_layout()
plt.show()

```



```
[183]: # Important features
important_features = ['partType', 'coolingRate', 'Nickel%', 'HeatTreatTime',
↳ 'Chromium%' , 'quenchTime']

# Subset the dataset to include only the important features
X_important = df[important_features]

# Split the dataset
X_train_imp, X_test_imp, y_train_imp, y_test_imp =
↳ train_test_split(X_important, df['Lifespan'], test_size=0.2, random_state=42)

# Initialize the Random Forest Regressor with the best parameters obtained from
↳ previous tuning
rf_model_imp = RandomForestRegressor(
    max_depth=15,
    n_estimators=387,
    random_state=42
)

# Train the model with reduced features
rf_model_imp.fit(X_train_imp, y_train_imp)

# Evaluate the model
y_pred_imp = rf_model_imp.predict(X_test_imp)
rmse_imp = root_mean_squared_error(y_test_imp, y_pred_imp)
r2_imp = r2_score(y_test_imp, y_pred_imp)
mae_imp = mean_absolute_error(y_test_imp, y_pred_imp)
msle_imp = mean_squared_log_error(y_test_imp, y_pred_imp)

print(f"Reduced Features RMSE: {rmse_imp:.2f}")
print(f"Reduced Features R2 Score: {r2_imp:.2f}")
print(f"Reduced Features MAE: {mae_imp:.2f}")
print(f"Reduced Features MSLE: {msle_imp:.2f}")
```

```
Reduced Features RMSE: 84.17
Reduced Features R2 Score: 0.93
Reduced Features MAE: 66.48
Reduced Features MSLE: 0.01
```